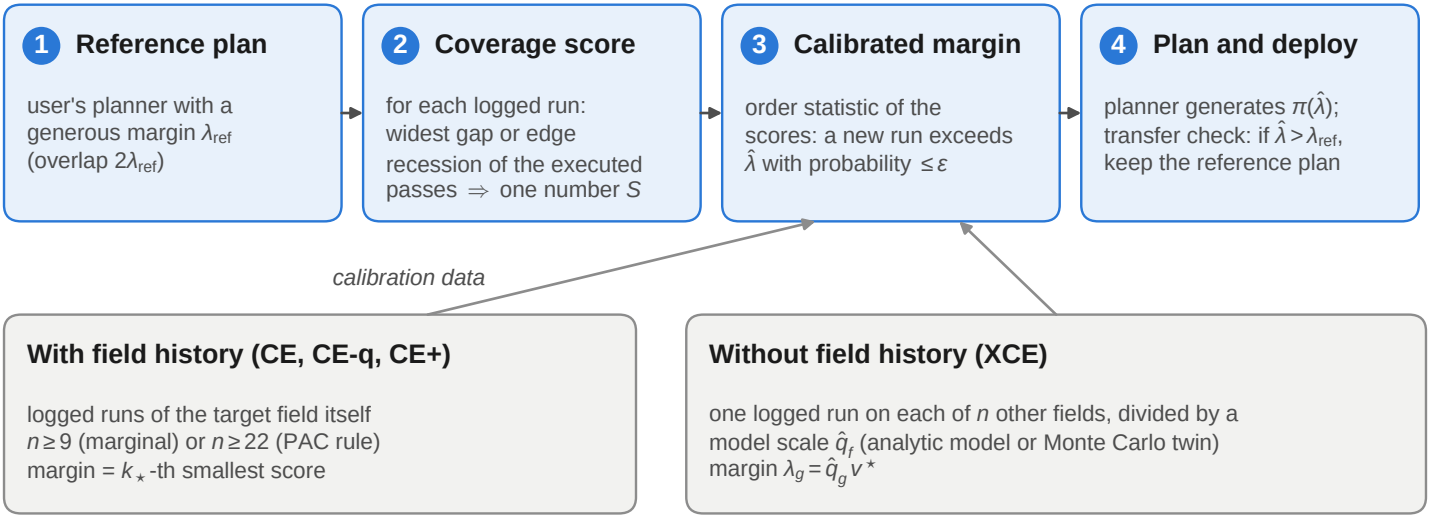
1 Reference plan
user's planner with a generous margin λ_{ref} (overlap $2\lambda_{ref}$)
2 Coverage score
for each logged run: widest gap or edge recession of the executed passes $\Rightarrow$ one number S
3 Calibrated margin
order statistic of the scores: a new run exceeds $\hat{\lambda}$ with probability $\leq \varepsilon$
4 Plan and deploy
planner generates $\pi(\hat{\lambda})$; transfer check: if $\hat{\lambda} > \lambda_{ref}$, keep the reference plan
calibration data
With field history (CE, CE-q, CE+)
logged runs of the target field itself
$n \geq 9$ (marginal) or $n \geq 22$ (PAC rule)
margin = $k_\star$-th smallest score
Without field history (XCE)
one logged run on each of n other fields, divided by a model scale $\hat{q}_f$ (analytic model or Monte Carlo twin)
margin $\lambda_g = \hat{q}_g v^\star$